# A COMMUNITY IN FORMATION: CORPUS CONSTRUCTION AND INITIAL ANALYSIS OF SLOVENE POSTS ON BLUESKY

Anonymous

Anonymous institution

This paper presents a corpus of 141,013 public Slovene-language posts by 432 authors on Bluesky, spanning August 2023 to mid-April 2026. The corpus was constructed through a protocol-native ATProto workflow combining relay-based discovery of Slovene-tagged posts, DID-resolved backfill of publicly available post histories from authors' Personal Data Servers, and post-level language filtering validated through manual annotation. As one of the first studies of Slovene-language use on Bluesky, the paper provides an initial descriptive account of an emerging public space on an open-protocol platform. The resulting corpus shows a marked increase in posting activity in late 2024. Reply posts account for 68.3% of the material, with a high reply share already visible before this increase, suggesting that conversational use predated the surge. Further observed patterns include sparse and largely account-concentrated hashtag use, as well as content sharing shaped primarily by Slovene news media, GIF-based informal exchange, and Bluesky-native tooling. More broadly, the paper shows how protocol-level access in ATProto can support the construction of a high-precision corpus for an under-resourced language without relying on platform-specific research APIs. At the same time, it highlights legal and ethical questions that arise when small and highly identifiable language communities become reconstructable from public decentralised infrastructures.

**Keywords:** Bluesky, ATProto, Slovene, corpus construction, social media

## 1 INTRODUCTION

Bluesky and the wider ATProto ecosystem have become increasingly relevant as a source of public social-media data (Failla & Rossetti, 2024; Quelle & Bovet, 2025). For smaller languages such as Slovene, this opens up an opportunity to document an emerging online community on a new, decentralised platform, and to consider how such documentation should be done. The paper is also motivated by the author's long-standing interest in Slovene online communities, familiarity with the emerging Slovene Bluesky space, and a broader interest in what open-protocol social media makes visible and researchable.

This paper describes an attempt to build a Slovene Bluesky corpus from public data. The main challenges are not only in collection, but in deciding what should count as a Slovene post. Language tags are useful but imperfect, and posts by Slovene-linked users are not automatically Slovene. The approach combines protocol-aware collection with explicit validation, but it still requires drawing a practical line around what counts as a Slovene post.

Four questions guide the work:

1. Can a usable Slovene Bluesky corpus be built from public ATProto data?
2. How far can collection move beyond a simple `langs=sl` strategy?
3. Which inclusion rules are reliable enough for a main corpus?
4. What basic properties are visible in the resulting corpus?

This paper makes two contributions. First, it presents a protocol-native workflow for constructing a high-precision corpus of Slovene Bluesky posts from public ATProto data, combining tagged discovery, multi-PDS backfill, and validated post-level language filtering. Second, it uses the resulting corpus for an initial descriptive analysis of temporal, interactional, and linking patterns in Slovene-language Bluesky posting.

## 2 RELATED WORK

### 2.1 Slovene language resources and social-media corpora

Written standard Slovene is well served by reference resources. Gigafida 2.0 (Krek et al., 2020a; Krek et al., 2019) is the main reference corpus, with about 1.1 billion words of written standard Slovene; ssj500k (Krek et al., 2020b) provides manually annotated training data; siParl 3.0 (Pančur et al., 2022) covers parliamentary proceedings from 1990 to 2022. The broader methodological context is surveyed in Gorjanc and Krek (2005).

The most directly relevant prior work is the JANES project (*Jezikovna analiza nestandardne slovenščine*), a central Slovene resource for user-generated content across social media and the web (Fišer et al., 2020). JANES mapped out much of the territory: the feasibility of building Slovene social-media corpora and the specific challenges they pose, including non-standard spelling and code-switching. The present work is much smaller in scale and ambition; it revisits some of the same questions in the context of Bluesky.

The closest language-specific Bluesky precedent located so far is Tjong Kim Sang (2025), who evaluates Bluesky as a successor data source for Dutch social-media

research after Twitter/X API restrictions made the TwiNL Twitter archive difficult to extend. That study collected 4.2 million Dutch Bluesky posts from November 2024 to May 2025 through keyword-based search and demonstrates compact corpus-linguistic analyses such as word-frequency comparison, daily rhythms, event response, hashtags, and political vocabulary. The present paper is narrower in language-community size, but differs methodologically by using protocol-level ATProto backfill from discovered authors and by making post-level language validation a central part of corpus construction.

### 2.2 Bluesky and ATProto

Prior work characterises Bluesky as both Twitter-like and decentralised, and links parts of its early growth to events around Twitter/X (Failla & Rossetti, 2024; Quelle & Bovet, 2025). Platform access conditions also matter here: Bluesky operated as an invite-only beta through 2023, regulating growth through waitlists and invite codes (Bluesky Social PBC, 2023b, 2023c), and only opened to the public on 6 February 2024, when invites were removed (Bluesky Social PBC, 2024). The present corpus cannot establish why individual Slovene users joined Bluesky, but the late-2024 surge visible in Section 5 should be interpreted against this broader platform context: public opening had already removed one barrier to entry, while the Slovene corpus shows its main surge only later, suggesting that public opening was a background enabling condition rather than a sufficient explanation on its own.

Technically, Bluesky is built on the AT Protocol (ATProto), a federated specification that separates identity from hosting: user identities are DIDs (decentralised identifiers), and public records are stored in per-account repositories on Personal Data Servers (Bluesky Social PBC, n.d.-a, 2023a; Kleppmann et al., 2024). ATProto also provides public synchronisation mechanisms: repository exports and a repository event stream, commonly called the firehose (Bluesky Social PBC, n.d.-b, n.d.-c). From a corpus-collection standpoint, this means that public posts can be collected through protocol-level mechanisms rather than through a separate academic access programme. DID resolution allows a collector to follow an account regardless of which PDS hosts it, a property exploited by the multi-PDS collection approach described in this paper.

Existing empirical work on Bluesky has mostly operated at large scale and focused on platform-wide, network-level, or topic-specific phenomena. Failla and Rossetti (2024) assembled a dataset of about 235 million posts from over four million users for behavioural analysis and reported large language-tag clusters, especially English, Japanese, and German. Quelle and Bovet (2025) used Bluesky data to characterise network topology, political polarisation, and user-generated feeds. More recent work

treats Bluesky's discovery mechanisms as research objects in their own right: Balduf et al. (2025) analyse starter packs as a mechanism for social bootstrapping, while Hanny et al. (2026) use multilingual keyword-based Bluesky collection for disaster analysis. These studies make clear that Bluesky is not simply a Twitter replacement: feeds, starter packs, repository access, and language metadata all affect what can be collected and interpreted.

Language-specific Bluesky corpus work appears sparse in the literature reviewed here. The Dutch study by Tjong Kim Sang (2025) is therefore a useful comparison point, while this paper instead offers a smaller but more validation-oriented case study for Slovene and for ATProto-based collection of a small language community.

### 2.3 Language identification in corpus construction

Building a language-specific corpus from a multilingual stream typically combines two signals: platform-supplied language metadata and automatic language identification. Neither is reliable on its own. Bluesky posts may contain a `langs` field with one or more BCP-47 language codes, and the official client auto-detects post languages while allowing users to override them (Bluesky PBC, n.d.-b). Large-scale Bluesky studies have also found that language metadata can be inconsistent or incomplete: Failla and Rossetti (2024) report non-standard and ill-formed language tags, while Hanny et al. (2026) caution that language metadata was inaccurate enough in a multilingual disaster crawl that filtering only by existing tags would risk missing relevant posts.

Automatic language identification tools—langid.py (Lui & Baldwin, 2012) and langdetect (Nakatani, 2010) are used here—are also imperfect, especially for short texts and closely related South Slavic languages. Using consensus across multiple detectors and requiring agreement with author-supplied tags improves precision, which is the rationale for the tiered inclusion strategy described in Section 3.

## 3 DATA COLLECTION AND METHODOLOGY

The corpus was built through a layered pipeline: tagged discovery through public relay mechanisms, seed-author identification, backfill from each author's home server, and post-level language filtering. Each phase is described below.[1]

[1] The collection scripts and supporting code are available at a repository anonymised for review. The corpus itself is not released as full text at this stage due to legal and ethical considerations; the intention is to provide an appropriate access format later (see Section 7).

### 3.1 Phase 1: Tagged discovery

Two complementary collection methods were used to discover Slovene-linked posts through public Bluesky relay infrastructure.

First, a **historical crawl** queried the Bluesky collection directory for public repositories containing `app.bsky.feed.post` records. The relevant relay endpoint is `listReposByCollection` (Bluesky PBC, n.d.-a; Bluesky Social PBC, n.d.-c). For each discovered repository, the author's DID was resolved to its home Personal Data Server (PDS), and the post records were fetched via `listRecords` from that PDS. Only posts whose `langs` field contained the tag `sl` or any `sl-*` subtag were retained.

Second, a **live firehose** subscriber connected to the public repository event stream (`subscribeRepos` on wss://bsky.network/xrpc) and processed post creation, update, and deletion events in real time, again retaining only Slovene-tagged posts (Bluesky Social PBC, n.d.-b). The firehose subscriber also handled account-level events (deletions, deactivations, suspensions) by marking the affected posts accordingly in the local store.

Both sources stored their results in local SQLite databases, with posts keyed by their AT-URI and soft-deletion tracking for posts that were later removed by their authors.

### 3.2 Phase 2: Seed author identification

Every unique author DID that appeared in at least one Slovene-tagged post from either discovery source was collected into a seed author list. The resulting list defines the author population for the rest of the pipeline. The corpus is therefore not a sample of the Bluesky post stream, but an author-centred collection of currently recoverable posts from a defined set of discovered Slovene-linked accounts.

### 3.3 Phase 3: Author-history backfill

For each seed author, the publicly available post history was retrieved through AT-Proto repository access. Each author's DID was resolved through a public directory: `plc.directory` for `did:plc:` identifiers, or the `.well-known/did.json` endpoint for `did:web:` identifiers. The resolved PDS URL was cached locally to avoid repeated lookups. Posts were then fetched via paginated `listRecords` calls to the author's home PDS.

Because resolution goes through the DID rather than a fixed domain, the backfill was not restricted to accounts on `bsky.social`; accounts hosted on other public

ATProto PDSs could be retrieved through the same DID-resolution workflow. This phase retrieved posts still publicly available at crawl time from each seed author, regardless of language tag, transforming the corpus from a tag-filtered post stream into an author-centred collection. Deleted posts that were no longer available at crawl time are not recoverable through this method and are out of scope.

The resulting expanded store contained 379,482 posts from 742 active authors, spanning the earliest surviving post (May 2023) through the collection date (April 2026).

### 3.4 Phase 4: Post-level language filtering

Each post in the expanded store was then classified individually to determine whether it should be included in the Slovene corpus. The classifier applied three binary signals:

- **Tag signal**: does the post's `langs` field contain `sl` or `sl-*`?
- **langid signal**: does the langid.py language identifier (Lui & Baldwin, 2012) assign the label `sl` to the post text?
- **langdetect signal**: does the langdetect library (Nakatani, 2010) assign Slovene a probability of at least 0.90?

Before classification, URLs and @-mentions were stripped from the post text. Posts with fewer than 20 alphabetic characters after stripping were treated separately, because language identification on very short texts is unreliable. The langdetect library was seeded deterministically (seed = 0) to ensure reproducible results.

For posts meeting the minimum length, the combination of signals determines the decision group. Table 1 summarises the mapping.

Table 1: Decision groups assigned by the language filtering step. Only groups marked as *included* are present in the final corpus.

| Decision group | Tag | langid | langdetect | Status |
|---|---|---|---|---|
| Core tag-supported* | yes | — | — | Included |
| Model consensus only | no | yes | yes | Included |
| langid only | no | yes | no | Included |
| langdetect only | no | no | yes | Included |
| Tag only | yes | no | no | Excluded |
| Short tagged | yes | *not evaluated* | | Excluded |

*Tag is *sl* and at least one of langid or langdetect agrees.

The core group (`core_tag_supported`) contains posts where the author tagged the post as Slovene *and* at least one language identifier agrees. The three identifier-supported review groups extend coverage to posts that were not tagged `sl` but are

classified as Slovene by the language identifiers. Posts tagged `sl` with no identifier support (`review_tag_only`) and posts too short for reliable classification (`review_short_tagged`) were excluded after manual validation confirmed their higher noise rate (see Section 4).

In addition to langid and langdetect, two further language identification tools—fastText (Joulin et al., 2017) and Lingua (Stahl, 2022)—were evaluated on the manually validated samples. On the comparable subset, all four detectors showed the same observed precision; the relevant difference was recall. The langid–langdetect pair recovered more acceptable Slovene posts than the alternatives, so the original detector pair was retained for the final corpus build.

### 3.5 Final corpus assembly

The core and the three included review groups were merged into a single final corpus, deduplicated by post URI. Posts from multiple collection runs (historical crawl, live firehose, and incremental updates) were combined in this step, with the first occurrence of each URI retained. The final corpus contains 141,013 posts by 432 authors; the remaining 310 authors in the expanded store had no posts that passed the language filter.

Each post in the analysis dataset is stored with its text, timestamp, author-supplied language tags, post-type flags (reply, quote), embed kind, and the per-post classifier outputs (decision group, langid label, langdetect probability). No linguistic annotation is included in the present corpus build.

### 3.6 Unit of analysis

The unit of analysis is the individual post. Language classification is applied per post, not per author or per conversation thread. A post in English by a Slovene-linked author inside an otherwise Slovene thread is classified on its own text and excluded if the text is not Slovene. This means the corpus represents Slovene-language posting, not Slovene-linked authorship in general.

## 4 VALIDATION

### 4.1 Precision validation

Validation was carried out through manual annotation of sampled posts. The strict core was evaluated using a random sample of 300 posts drawn uniformly from the

`core_tag_supported` group. All 300 sampled posts were judged acceptable as Slovene, yielding an estimated precision of 100% on the core layer.

The broader review layer was evaluated with a separate random sample of 300 posts to determine which decision groups were clean enough to include. A breakdown by decision group showed that noise was concentrated in the tag-only and short-tagged subgroups. Within this 300-post sample, every row drawn from the three model-supported groups listed in Table 1 was judged acceptable. The overall sample-level acceptability rate of 72% (214 `Slovene-dominant`, 2 `Mixed-with-Slovene`, 40 `Not-Slovene`, 44 `Undeterminable/too-short`) is driven by the two noisier groups and therefore does not reflect the quality of the three model-supported groups that were ultimately included in the corpus.

Additional dedicated follow-up checks on the two smaller model-supported groups confirmed this: in the `review_langid_only` sample of 50 posts, 49 were labelled `Slovene-dominant` and 1 `Mixed-with-Slovene`; in the `review_langdetect_only` sample of 50 posts, 49 were labelled `Slovene-dominant` and 1 `Not-Slovene`.

These results support a final corpus that includes the strict core and all three model-supported review groups, while excluding the tag-only and short-tagged material whose noise the validation confirmed.

### 4.2 False-negative estimation

Precision validation confirms that posts included in the corpus are Slovene. A separate analysis estimated the false-negative rate, i.e. the fraction of Slovene posts from known authors that were excluded by the pipeline. A random sample of 200 posts from seed authors that (a) were not included in the final corpus, (b) carried no Slovene language tag, and (c) met the project's minimum-text rule of 20 alphabetic characters after URL and @-mention stripping was manually annotated. Of these, 17 posts (8.5%) were judged to be Slovene-dominant or Mixed-with-Slovene.

A rough extrapolation to the full pool of 97,968 excluded untagged posts from known authors that met the same minimum-text rule suggests that roughly 8,300 Slovene posts from the same author set were not captured by the current validated rule. This corresponds to approximately 5.9% of the size of the final corpus. The corpus should therefore be described as a high-precision collection of Slovene posts from discovered Slovene-linked authors under the current rule, not as a complete census of all Slovene writing on Bluesky.

## 5 EMPIRICAL ANALYSIS

### 5.1 Corpus profile

The final corpus contains 141,013 posts by 432 unique authors, spanning 9 August 2023 to 16 April 2026. Table 2 summarises its main properties. Posts are distributed across four decision groups: 85,094 `core_tag_supported` (60.3%), 46,134 `review_model_consensus_only` (32.7%), 5,632 `review_langid_only` (4.0%), and 4,153 `review_langdetect_only` (2.9%).

Table 2: Summary statistics for the final Slovene Bluesky corpus.

| Measure | Value |
|---|---|
| Posts | 141,013 |
| Unique authors | 432 |
| Date range | 9 Aug 2023–16 Apr 2026 |
| Median post length | 86 characters |
| Replies | 96,278 (68.3%) |
| Original posts | 41,517 (29.4%) |
| Non-reply quote posts | 3,218 (2.3%) |
| Authors for 50% of posts | 22 |

### 5.2 Growth and participation over time

Monthly metadata were used to distinguish raw post growth from the formation of a more stable author community. For each calendar month, the analysis counted total posts, active authors, first-seen authors, returning authors, cumulative authors, posts per active author, and simple one- and three-month first-seen author retention. All timestamps are treated as UTC. The final month, April 2026, is partial and is therefore excluded from retention windows, monthly period medians, and the statistical comparisons below.

One break stands out clearly in the monthly series. In November 2024, Slovene post volume increased from 31 in October 2024 to 5,984 the following month, while the number of active authors increased from 13 to 196. The number of first-seen authors rose from 8 in October to 172 in November. By the end of November 2024, the cumulative number of observed authors had reached 202, and by the end of December it had reached 252. Figure 1 shows the full monthly post curve.

This break is visible not only in volume but also in participation. Through October 2024, the corpus contains 178 posts by 30 authors, with a median of 10 posts and 3 active authors per complete month. From November 2024 onward, the corpus

contains 140,835 posts by 429 authors, with median complete-month values of 7,958 posts and 210 active authors. Median posts per active author also rises from 2.33 before the break to 37.73 after it, indicating that the change is not only an increase in the number of participating authors. Table 3 summarises the period comparison.

Table 3: Growth and participation metrics before and after the late-2024 surge. Monthly medians exclude the partial final month, April 2026.

| Measure | Through Oct 2024 | Nov 2024 onward |
|---|---:|---:|
| Total observed posts | 178 | 140,835 |
| Total active authors | 30 | 429 |
| Median monthly posts | 10 | 7,958 |
| Median monthly active authors | 3 | 210 |
| Median posts per active author | 2.33 | 37.73 |

Retention by first-seen month points in the same direction: the late-2024 surge does not appear to consist of one-off activity alone. Of the 172 authors first observed in November 2024, 86.6% posted again in December 2024 and 94.2% posted at least once in the following three months. Of the 50 authors first seen in December 2024, 72.0% posted again in January 2025 and 82.0% posted at least once in the following three months. While these figures do not explain why authors became active, they do suggest that most continued to participate beyond their first appearance.

The late-2024 break also occurs well after Bluesky's public opening in February 2024 (Bluesky Social PBC, 2024), indicating that the removal of invite-only access was a necessary background condition for growth but not, by itself, sufficient to explain the much later Slovene surge. The timing furthermore coincides with a wider post-election migration wave to Bluesky following the November 2024 U.S. presidential election (Parvini, 2024). These events are treated here as contextual reference points rather than causal explanations.

Taken together, the observed changes represent a large and systematic shift across monthly post counts, numbers of active authors, and posts per active author.[2]

### 5.3 Temporal patterns

Beyond growth and participation, temporal regularities in posting activity offer another perspective on how this community uses the platform.

[2] Mann–Whitney U tests on complete pre- and post-November 2024 months confirm this pattern statistically, with Cliff's delta = 1.00 and two-sided $p \approx 1.6 \times 10^{-6}$ in each case, used descriptively rather than causally.

Figure 1: Monthly Slovene post counts (bars, left axis) and active authors (line, right axis), August 2023–April 2026. Dashed lines: Bluesky's public opening (6 Feb 2024) and the post-election migration wave (Nov 2024). Dotted lines: EU Parliament elections (9 Jun 2024) and Slovenian parliamentary elections (22 Mar 2026). April 2026 bar is partial (ends 16 April).

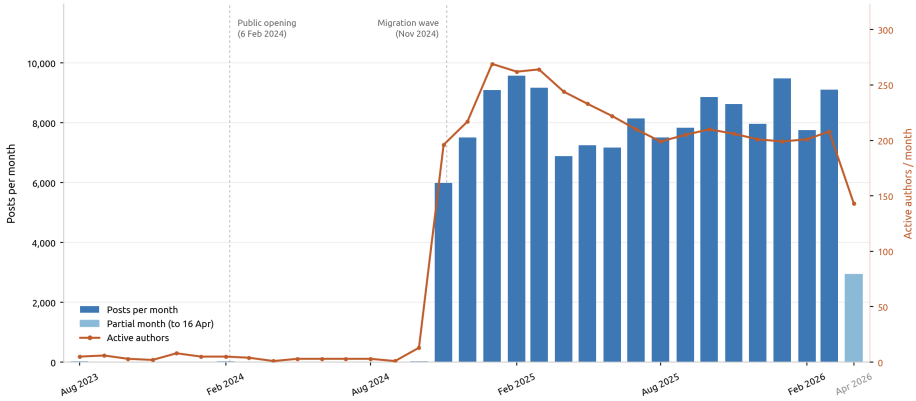


Posting activity follows a clear daily rhythm consistent with the Central European time zone. A morning peak occurs at 09:00–10:00 UTC and a stronger evening peak at 18:00–19:00 UTC, corresponding roughly to late morning and evening in Slovenia depending on daylight saving time. Activity is lowest between 01:00 and 04:00 UTC. This pattern is consistent with a community of primarily Slovenian-timezone users.

### 5.4 Interactional structure

Using a mutually exclusive hierarchy that counts reply-quote combinations as replies, the corpus is dominated by reply posts: 96,278 of 141,013 included posts (68.3%) are replies, while 41,517 (29.4%) are original posts and 3,218 (2.3%) are non-reply quote posts. The median reply is 77 characters, compared to 114 characters for original posts. In this corpus, Slovene Bluesky use is therefore strongly conversational.

This pattern is not only a result of the late-2024 growth surge. Before November 2024, replies already accounted for 54.5% of posts; from November 2024 onward, their share rises to 68.3%. The post-surge influx of authors is therefore consistent with a continuation and strengthening of an interactional pattern that was already present, rather than the introduction of a fundamentally different one.

The author-level distribution tells a similar story. Among the 246 authors with at least 20 posts in the final corpus, the median reply rate is 0.66 (IQR 0.38–0.84), and only 9 authors (3.7%) have a reply rate of zero. At the same time, 74 authors (30.1%) have reply rates above 0.8. This suggests that the dominance of replies is not only an

aggregate artefact of corpus-level totals, but a feature that extends beyond a small number of highly active accounts. A simple robustness check leads to the same conclusion: the median monthly reply share is nearly identical in the full corpus and in the subset of authors with at least 20 posts (67.1% vs. 67.3%), indicating that the conversational character of the corpus is not driven mainly by marginal one-off accounts.

A small broadcast-oriented minority is nevertheless visible. Using simple descriptive thresholds based on reply rate, 39 authors in the subset of authors with at least 20 posts fall into a broadcast-heavy group (reply rate below 0.2), 133 into a mixed group, and 74 into a strongly conversational group (reply rate above 0.8). These labels are only descriptive heuristics, not validated social categories, but they help illustrate the broad split already visible in the data: a limited number of accounts contribute mostly non-reply content, while a substantial share of active authors use the platform primarily for exchange and ongoing dialogue.

Posts with emoji occur at a rate of 19.5%, with at least one external URL in the post text at 12.6%, hashtags at 4.9%, and @-mentions at 1.2%. The low hashtag rate is consistent with the conversational character of the corpus. Across the whole corpus, 11.2% of posts embed a structured external link card, 8.0% include images, 2.3% are non-reply quote posts embedding another post, 0.5% contain video, and 0.2% combine a quote-post with attached media. Video is notably rare.

### 5.5 Author concentration and infrastructure

Posting activity is heavily concentrated. As Table 2 shows, just 22 authors are responsible for half of the corpus. The top 10% of authors (43 users) account for 69.6% of all posts, while the median author contributed 35 posts over the entire collection period. This skewed distribution is characteristic of social media corpora generally, though the concentration here is particularly pronounced given the small and still-forming community.

Regarding platform infrastructure, 423 of 432 authors (97.9%) are hosted on Bluesky's own PDS infrastructure (`*.host.bsky.network`), while 9 (2.1%) are hosted on `eurosky.social`, an independent PDS. This confirms that the corpus is not restricted to a single server and that DID-based PDS resolution is operational in the collection workflow.

Handle choice is slightly more varied. Of the 432 authors, 395 (91.4%) use a `*.bsky.social` handle, 30 (6.9%) use a user-owned custom domain, and 7 (1.6%) use a `*.eurosky.social` subdomain. The custom-domain group contributed 13,303

posts (9.4% of the corpus). Two authors combine a custom domain handle with `eurosky.social` hosting.

Together, these figures show that most corpus authors use the default Bluesky configuration, while a smaller minority changed either their public handle, their hosting provider, or both.

### 5.6 Hashtag and linking practices

Hashtag use is sparse throughout the corpus. Only 6,879 of 141,013 posts (4.9%) contain at least one hashtag, suggesting that broad hashtag coordination is not central to this corpus. The most frequent hashtags are often highly concentrated: among the 20 most frequent hashtags, 18 have a single author contributing more than half of all uses. Several of the top hashtags, including `#supslovenija`, `#supdeska`, and `#aquamarina`, are used exclusively by a single account and therefore function more as repeated self-tagging than as markers of a wider community topic. At the same time, a small minority of hashtags shows more distributed authorship, such as `#pohvalanadan` (91 occurrences, 28 authors) and `#ema2025` (67 occurrences, 15 authors), suggesting that wider circulation does occur but is not the dominant pattern.

External link sharing is more prevalent than hashtagging. A total of 17,838 posts (12.6%) contain at least one external domain. The most frequently linked domain is `media.tenor.com` (2,613 occurrences across 121 authors), indicating that GIF-based informal exchange is a prominent and broadly shared interactional practice. Other highly visible domains are more concentrated: `ebx.sh`, a Bluesky-related link-sharing utility, appears 1,804 times but is used by only two authors, while `tehnozvezdje.si` is likewise concentrated in a small number of accounts. These domains point to two different practices: broadly shared informal exchange on the one hand, and repeated account-specific linking on the other.

At the level of domain categories, linked content is shaped primarily by Slovene news media (35.8% of domain occurrences), followed by GIF and media-sharing services (14.5%), platform tooling and Bluesky utilities (11.2%), and video platforms (8.0%). The news category is distributed across several outlets rather than concentrated in a single source, pointing towards recurring current-affairs engagement. At the same time, the prominence of GIF sharing and Bluesky-native tools suggests a space shaped not only by information sharing, but also by informal exchange and platform-specific habits.

Link-sharing also differs sharply by post type. External links occur in 32.2% of original posts but in only 4.5% of replies. This is consistent with the broader interactional pat-

tern observed in Section 5.4: original posts are comparatively more likely to introduce or circulate content, while replies function primarily as conversational turns. Taken together, the hashtag and domain patterns suggest a small conversational public space oriented around informal exchange and recurring news-media engagement, rather than one organised primarily through hashtag-driven topic publics.

### 5.7 Linguistic profile and multilinguality

A notable feature of the corpus is the frequency of mixed language tagging. Of the 141,013 posts, 58,416 (41.4%) were tagged with both `sl` and `en` language codes. By contrast, langid detected a non-Slovene primary language in only 7.3% of posts. The predominant non-Slovene labels assigned by langid were Croatian (*hr*: 4,506 posts) and Bosnian (*bs*: 2,285 posts); closely related South Slavic languages for which cross-label confusion by automatic detectors is common. langid assigned English to only 278 posts.

The high rate of `sl+en` tagging without corresponding English detection by langid suggests that mixed language metadata often reflects client-side tagging behaviour rather than sustained English-dominant content. The corpus nevertheless contains English lexical insertions within otherwise Slovene posts, which are retained as authentic examples of mixed-register writing.

## 6 DISCUSSION

The empirical findings in the previous section can be read in two ways: as descriptive statistics about a dataset, or as a portrait of a community caught in the process of forming. This section takes the second reading.

### 6.1 Early activity and community formation

The late-2024 break visible in the corpus can be read not only as a change in volume, but also as part of the early formation of a Slovene Bluesky public during a broader period of platform change and migration. Early activity on newly adopted platforms often includes meta-platform orientation: posts about the platform itself, comparisons with older platforms, and negotiation of emerging features, norms, and expectations. Similar dynamics were also visible in public discussion around platforms such as Mastodon and Threads, where users likewise frequently compared them to Twitter/X; a comparable orientation is plausible here as well, even if it was not directly measured in the present study. From a user perspective, a platform arguably becomes fully

adopted only once this phase fades into the background and ordinary use takes over: when users stop foregrounding how the platform works and simply use it. The late-2024 surge and the retention patterns observed in Section 5 suggest that, for many authors, Bluesky was by that point becoming a more regular space of participation rather than remaining only a novelty. Whether this transition is also visible in topic choice, register, or interactional tone remains a question for future work.

### 6.2 Dialogue over broadcasting

The corpus is dominated by replies, and this pattern is not confined to a small number of marginal or highly active accounts. In that sense, Slovene Bluesky use appears to be oriented more towards exchange than towards one-way broadcasting. Because reply-dominant use is already visible before the late-2024 surge, the post-surge expansion is better understood as a continuation and strengthening of an existing interactional tendency than as a shift in kind.

### 6.3 Small but concentrated

The Slovene Bluesky space captured here is small and highly concentrated, with a limited core of authors contributing a large share of the corpus. Its communicative profile is shaped less by broad hashtag coordination than by conversational exchange, recurring news-media linking, and widely shared informal practices such as GIF use. Taken together, these patterns suggest a small public oriented more around interaction and repeated content-sharing than around large-scale topic publics. At the same time, the size of the corpus limits stronger claims about internal subgroup structure.

### 6.4 ATProto as a research infrastructure

From a methodological standpoint, the experience of building this corpus confirms that ATProto's open design has real consequences for research. Public synchronisation and repository mechanisms made a backfill of the discovered Slovene-linked authors achievable without special research access (Bluesky Social PBC, n.d.-b, n.d.-c). For researchers working on small or under-resourced languages, where corpora are scarce and institutional resources are limited, this openness matters. The Slovene community on Bluesky may be small, but the corpus captures much of the community's publicly available posting history.

## 7 LIMITATIONS AND ETHICS

### 7.1 Coverage limits

The corpus should not be described as a complete census of all Slovene Bluesky posts ever published. It covers currently recoverable public posts from discovered Slovene-linked accounts; deleted posts are excluded by design. The false-negative analysis (Section 4) estimates an 8.5% miss rate within the pool of excluded untagged posts from known authors that met the project's minimum-text rule; the true miss rate across the full user population is unknown and likely higher.

The methodology uses post-level rather than thread-level language classification. Slovene-linked users may post in English or other languages within otherwise Slovene conversations, and such replies are not automatically treated as Slovene corpus material.

### 7.2 Open data and reproducibility

The AT Protocol's openness is what made this corpus possible. The collection described here used public protocol mechanisms rather than a paid or application-only research API. This does not mean that access is unlimited: Bluesky documents rate limits for its services, and the ATProto sync specification explicitly allows service providers to apply rate limits and data quotas (Bluesky PBC, n.d.-c; Bluesky Social PBC, n.d.-c). For a small-language community like Slovene, however, the availability of public synchronisation mechanisms is still a practically valuable property, though it also raises specific questions about redistribution.

### 7.3 Privacy and identifiability

The same openness that makes collection easy raises a concern worth stating plainly. The Slovene Bluesky community, at the time of collection, consists of 432 identified authors; a small number. In most cases, the individuals behind the accounts can be identified with little effort: handles are publicly visible, many users link their real names or professional identities through custom domains, and the posting history of any individual is freely accessible through DID resolution. The corpus therefore represents not just a dataset but a detailed longitudinal record of specific, identifiable people.

This is not unique to Bluesky or to this corpus; it is a structural feature of public social-media research. *Public* does not mean *anonymous*: social-media content can remain

identifiable even when obvious identifiers are removed, and direct quotation or full-text redistribution can make reverse identification easier (Ayers et al., 2018; Franzke et al., 2019). Users who post under their real names may not have anticipated that their posting history would be compiled into a citable research corpus. Following standard practice, full post texts are not reproduced in this article, and any future redistribution should take a conservative form to limit re-identification risk.

### 7.4 Corpus redistribution on a decentralised platform

Whether and how to redistribute this corpus is not straightforward, and this paper does not attempt to resolve the question. ATProto supports content and account deletion, and its synchronisation specification explicitly warns against public redistribution of static repository snapshots in bulk, as mirrors are expected to respect record deletions and account-status changes (Bluesky Social PBC, n.d.-c). A static corpus distributed as a file cannot fully reflect this behaviour: posts deleted on Bluesky remain present in copies that have already been downloaded elsewhere. From a European data protection perspective, public posts may still constitute personal data when authors are identifiable (European Parliament and Council of the European Union, 2016).

Identifier-based redistribution has been used in Twitter research to balance reproducibility with deletion and access constraints (Kinder-Kurlanda et al., 2017). For ATProto, an equivalent and widely accepted redistribution model has not yet emerged. In the present study, the corpus is therefore described but not released as a full-text dataset. The methodology and collection code are made available, and it is in principle possible to reassemble the corpus from public ATProto data under the conditions described in this paper.

At the same time, it is the author's intention to make the corpus available in a form that is legally and ethically defensible. A carefully defined redistribution model would enable reuse without requiring individual researchers to repeat the full collection process on their own. However, determining an appropriate access format for a small and identifiable language community requires further consideration and is treated as future work rather than resolved here.

## 8 CONCLUSION

This paper has described a workflow for building a high-precision corpus of public Slovene Bluesky posts and used that corpus for an initial look at the interactional and temporal structure of Slovene activity on the platform.

The pipeline combines relay-based tagged discovery (historical crawl and live firehose), DID-based backfill from individual Personal Data Servers, and post-level language filtering validated through manual annotation. It is fully protocol-native: no special API keys, no institutional data-access agreements, and no restriction to a single hosting server. For small-language communities where research infrastructure is limited, this openness is a practical advantage.

The corpus of 141,013 posts by 432 authors shows a sharp late-2024 expansion in recoverable activity, strong conversational use, and highly concentrated participation. Sparse hashtag use and link-sharing shaped largely by Slovene news media, GIF-based exchange, and platform-native tools further suggest a small public marked more by recurring interaction than by broad topic coordination.

This corpus should be treated as a snapshot of currently recoverable public posts from discovered Slovene-linked authors, not a complete census of all Slovene writing on Bluesky. Deleted posts are excluded by design, and within the pool of excluded untagged posts from known authors that meet the project's minimum-text rule, the false-negative rate is estimated at approximately 8.5%. The main next step is not only technical enrichment, but a redistribution model that makes the corpus useful without ignoring the rights and expectations of the people whose posts make it possible.

## 9 FUTURE WORK

The most immediate next step is a more reusable access format for the corpus — one that makes it possible to work with the data without re-running the entire pipeline, while remaining consistent with the legal and ethical constraints discussed in Section 7.

A second direction is linguistic enrichment. Adding annotation would make the corpus more useful for further work on non-standard Slovene, code-switching, and interaction in short-form social-media text. It would also open the way to content-oriented analyses that go beyond the descriptive focus of the present paper, including topic analysis, studies of register, lexical choice, and discourse patterns in annotated texts. Given the high share of replies in the corpus, thread-level analysis would be especially relevant for examining how interaction unfolds conversationally rather than only at the level of isolated posts.

A third direction is comparative and longitudinal research. The present corpus could be extended forward in time in order to follow later developments in the Slovene Bluesky space, including possible shifts in activity, participation, and language use.

It would also be valuable to compare this corpus with other Slovene online communities and platforms, such as Twitter/X, Mastodon, or Threads, in order to examine how platform architecture, migration context, and local uptake shape interactional patterns, linking practices, and public conversation. More broadly, the same collection approach could be adapted to other small-language communities on Bluesky, supporting corpus construction for smaller language communities on emerging social-media platforms.


## ACKNOWLEDGEMENTS

Generative AI tools were used during the coding, implementation, analysis support, and writing phases of this paper, primarily ChatGPT/Codex (GPT-5.4) and Claude Code (Claude Sonnet 4.6). The research topic, overall direction of the work, and final interpretive decisions were determined by the author. All generated material was reviewed, revised where necessary, and finalised by the author.


## REFERENCES


Ayers, J. W., Caputi, T. L., Nebeker, C., & Dredze, M. (2018). Don't quote me: Reverse identification of research participants in social media studies. *NPJ Digital Medicine*, *1*(1), 30. https://doi.org/10.1038/s41746-018-0036-2

Balduf, L., Sokoto, S., Baronchelli, A., Castro, I., Król, M., Tyson, G., Pavlou, G., Scheuermann, B., & Ascigil, O. (2025). Bootstrapping social networks: Lessons from Bluesky starter packs. *Proceedings of the Nineteenth International AAAI Conference on Web and Social Media*, *19*, 178–192. https://doi.org/10.1609/icwsm.v19i1.35810

Bluesky PBC. (n.d.-a). *com.atproto.sync.listReposByCollection*. Retrieved April 21, 2026, from https://docs.bsky.app/docs/api/com-atproto-sync-list-repos-by-collection

Bluesky PBC. (n.d.-b). *Posts*. Retrieved April 21, 2026, from https://docs.bsky.app/docs/advanced-guides/posts

Bluesky PBC. (n.d.-c). *Rate limits*. Retrieved April 21, 2026, from https://docs.bsky.app/docs/advanced-guides/rate-limits

Bluesky Social PBC. (n.d.-a). *Repository*. Retrieved April 21, 2026, from https://atproto.com/specs/repository

Bluesky Social PBC. (n.d.-b). *Streaming data*. Retrieved April 21, 2026, from https://atproto.com/guides/streaming-data

Bluesky Social PBC. (n.d.-c). *Sync*. Retrieved April 21, 2026, from https://atproto.com/specs/sync

Bluesky Social PBC. (2023a). The AT Protocol [Technical specification]. Retrieved April 1, 2026, from https://atproto.com/guides/overview

Bluesky Social PBC. (2023b). *Bluesky private beta update*. Retrieved April 22, 2026, from https://bsky.social/about/blog/3-2-2023-bluesky-beta-app

Bluesky Social PBC. (2023c). *Private beta update & roadmap*. Retrieved April 22, 2026, from https://bsky.social/about/blog/6-02-2023-beta-update

Bluesky Social PBC. (2024). *Join Bluesky today (bye, invites!)* Retrieved April 22, 2026, from https://bsky.social/about/blog/02-06-2024-join-bluesky

European Parliament and Council of the European Union. (2016). Regulation (EU) 2016/679 of the European Parliament and of the Council of 27 April 2016 on the protection of natural persons with regard to the processing of personal data and on the free movement of such data, and repealing Directive 95/46/EC (General Data Protection Regulation). Retrieved April 21, 2026, from https://eur-lex.europa.eu/eli/reg/2016/679/oj/eng

Failla, A., & Rossetti, G. (2024). "I'm in the Bluesky Tonight": Insights from a year worth of social data. *PLOS ONE*, *19*(11), e0310330. https://doi.org/10.1371/journal.pone.0310330

Fišer, D., Ljubešić, N., & Erjavec, T. (2020). The JANES project: Language resources and tools for Slovene user generated content. *Language Resources and Evaluation*, *54*, 223–246. https://doi.org/10.1007/s10579-018-9425-z

Franzke, A. S., Bechmann, A., Zimmer, M., Ess, C., & Association of Internet Researchers. (2019). Internet research: Ethical guidelines 3.0. Retrieved April 21, 2026, from https://aoir.org/reports/ethics3.pdf

Gorjanc, V., & Krek, S. (Eds.). (2005). *Študije o korpusnem jezikoslovju*.

Hanny, D., Schmidt, S., & Resch, B. (2026). Bluesky as a social media data source for disaster management: Investigating spatio-temporal, semantic and emotional patterns for floods and wildfires. *Journal of Computational Social Science*, *9*, 16. https://doi.org/10.1007/s42001-025-00448-x

Joulin, A., Grave, E., Bojanowski, P., & Mikolov, T. (2017). *fastText language identification* [Language identification model; based on: Joulin et al. (2017) Bag of Tricks for Efficient Text Classification, EACL 2017, https://aclanthology.org/E17-2068/]. Retrieved April 23, 2026, from https://fasttext.cc/docs/en/language-identification.html

Kinder-Kurlanda, K., Weller, K., Zenk-Möltgen, W., Pfeffer, J., & Morstatter, F. (2017). Archiving information from geotagged tweets to promote reproducibility and comparability in social media research. *Big Data & Society*, *4*(2). https://doi.org/10.1177/2053951717736336

Kleppmann, M., Frazee, P., Gold, J., Graber, J., Holmgren, D., Ivy, D., Johnson, J., Newbold, B., & Volpert, J. (2024). Bluesky and the AT Protocol: Usable decentralized social media. *Proceedings of the ACM CoNEXT-2024 Workshop on the Decentralization of the Internet*, 1–7. https://doi.org/10.1145/3694809.3700740

Krek, S., Arhar Holdt, Š., Erjavec, T., Čibej, J., Repar, A., Gantar, P., Ljubešić, N., Kosem, I., & Dobrovoljc, K. (2020a). Gigafida 2.0: The reference corpus of written standard Slovene. In N. Calzolari, F. Béchet, P. Blache, K. Choukri, C. Cieri, T. Declerck, S. Goggi, H. Isahara, B. Maegaard, J. Mariani, H. Mazo, A. Moreno, J. Odijk, & S. Piperidis (Eds.), *Proceedings*

*of the twelfth language resources and evaluation conference* (pp. 3340–3345). European Language Resources Association. https://aclanthology.org/2020.lrec-1.409/

Krek, S., Erjavec, T., Dobrovoljc, K., Gantar, P., Arhar Holdt, Š., Čibej, J., & Brank, J. (2020b). The ssj500k training corpus for Slovene language processing. In D. Fišer & T. Erjavec (Eds.), *Jezikovne tehnologije in digitalna humanistika: Zbornik konference* (pp. 24–33). Inštitut za novejšo zgodovino. https://doi.org/10.5281/zenodo.13954583

Krek, S., Erjavec, T., Repar, A., Čibej, J., Arhar Holdt, Š., Gantar, P., Kosem, I., Robnik Šikonja, M., Ljubešić, N., Dobrovoljc, K., Laskowski, C., Grčar, M., Holozan, P., Šuster, S., Gorjanc, V., Stabej, M., & Logar, N. (2019). *Corpus of written standard Slovene Gigafida 2.0*. Slovenian language resource repository CLARIN.SI. http://hdl.handle.net/11356/1320

Lui, M., & Baldwin, T. (2012). Langid.py: An off-the-shelf language identification tool. *Proceedings of the ACL 2012 System Demonstrations*, 25–30. https://aclanthology.org/P12-3005/

Nakatani, S. (2010). Language detection library for Java [Software library]. https://github.com/shuyo/language-detection

Pančur, A., Erjavec, T., Meden, K., Ojsteršek, M., Šorn, M., & Blaj Hribar, N. (2022). *Slovenian parliamentary corpus (1990-2022) siParl 3.0*. Slovenian language resource repository CLARIN.SI. http://hdl.handle.net/11356/1748

Parvini, S. (2024). *Bluesky has added 1 million users since the US election as people seek alternatives to X*. Retrieved April 22, 2026, from https://apnews.com/article/d78de194e0e85f803652b70b161a963b

Quelle, D., & Bovet, A. (2025). Bluesky: Network topology, polarization, and algorithmic curation. *PLOS ONE*, *20*(2), e0318034. https://doi.org/10.1371/journal.pone.0318034

Stahl, P. M. (2022). Lingua: The most accurate natural language detection library for Python [Software library]. https://github.com/pemistahl/lingua-py

Tjong Kim Sang, E. (2025). Using Bluesky for social media analysis. *Digital Humanities Benelux 2025*. https://doi.org/10.5281/zenodo.15650889

# SKUPNOST V NASTAJANJU: GRADNJA KORPUSA IN ZAČETNA ANALIZA SLOVENSKIH OBJAV NA PLATFORMI BLUESKY

Prispevek predstavlja korpus 141.013 javnih objav v slovenščini, ki jih je na mikroblogerski platformi Bluesky med avgustom 2023 in sredino aprila 2026 objavilo 432 avtorjev. Korpus je bil oblikovan z večstopenjskim postopkom na ravni protokola AT, ki zajema odkrivanje objav z oznako *sl*, rekonstrukcijo javno dostopnih zgodovin objav posameznih računov ter filtriranje posameznih objav glede na jezik. Kot ena prvih empiričnih, jezikovno specifičnih študij rabe platforme Bluesky za slovensko jezikovno skupnost se prispevek osredotoča na zgodnje faze oblikovanja novega, na odprtem protokolu temelječega javnega prostora. Zbrani podatki nakazujejo izrazit porast aktivnosti konec leta 2024. Odgovori na objave predstavljajo 68,3 % vseh objav, pri čemer je njihov delež visok že pred porastom aktivnosti. Opazna je tudi redka in avtorsko koncentrirana raba ključnikov, večinoma s slovenskimi mediji povezano deljenje novic ter pogosta raba GIF-ov, kar kaže na neformalnost komunikacije. Prispevek pokaže, kako lahko dostop na ravni protokola AT omogoča gradnjo korpusa za jezike z manjšim številom govorcev, hkrati pa odpira pravna in etična vprašanja, povezana z dostopnostjo javnih podatkov manjših jezikovnih skupnosti.

*Opomba k jeziku: Angleščina je v tem prispevku uporabljena izključno z namenom nagovarjanja širše mednarodne raziskovalne skupnosti na področjih digitalne humanistike, obdelave naravnega jezika in ATProta. Uporaba angleščine nikakor ne zmanjšuje pomena slovenščine v raziskovalnem in spletnem prostoru, katere vlogo še naprej dosledno in odločno zagovarjam.*

**Ključne besede:** Bluesky, ATProto, slovenščina, korpus, družbena omrežja

To delo je ponujeno pod licenco Creative Commons: Priznanje avtorstva-Deljenje pod enakimi pogoji 4.0 Mednarodna.

This work is licensed under the Creative Commons Attribution-ShareAlike 4.0 International.

https://creativecommons.org/licenses/by-sa/4.0/